

Transaction Concepts through SQL

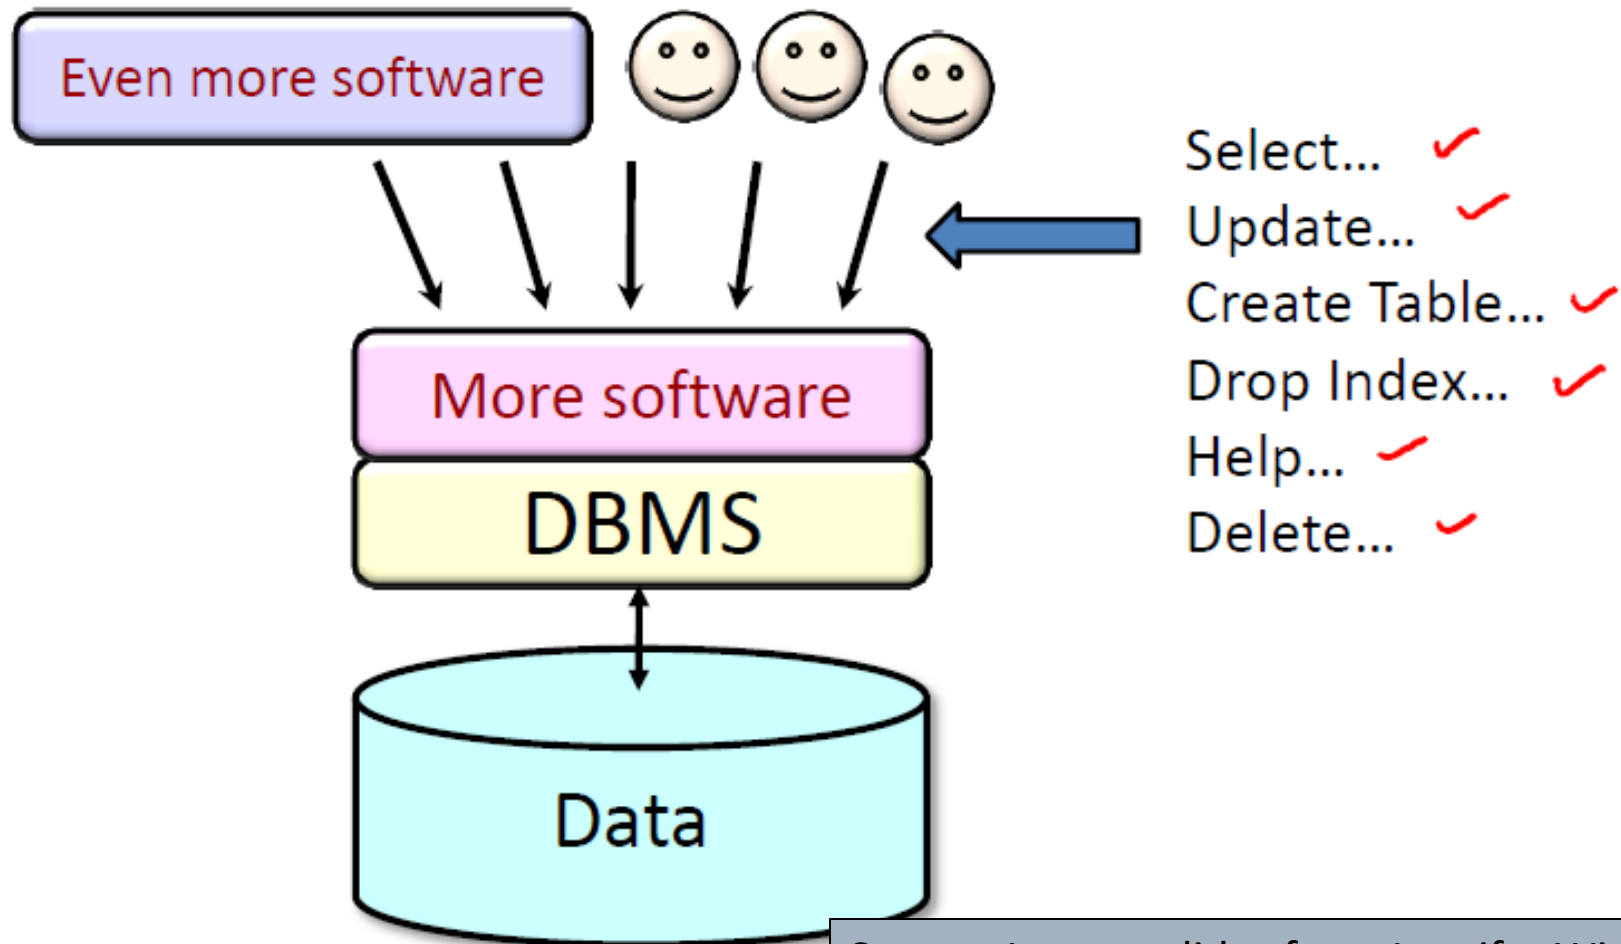


PM Jat
pm_jat@daiict.ac.in

Why do we study Transactions

- Study of Transactions addresses following two requirements –
 - Enable concurrent access of databases - multiple users need to concurrently access and update databases.
 - Resilience to system failures - While databases are updated system might crash in between.

Concurrent Database Access



Source: Lecture slides from Jennifer Widom

Execution of SQL statements

- SQL statements higher level expressions, and their execution involves execution of

```
UPDATE employee  
SET salary = salary + 3000 WHERE ssn = '1234';
```

- Let us say that transactions are executed in a sequence of logical operations - Read X, Modify X, Write X.

Execution of SQL statements

UPDATE employee

SET salary = salary + 3000 WHERE ssn = '1234';

- Using the said model, let us say that above SQL statements is executed as –
 - Read Salary (let us assume it is 50,000),
 - Modify it to 53000 (by adding 3000)
 - Write 53000 to database

- At the end of transaction, database has 53000 as value of “the salary”

Relations used here

- ☐ Employee(SSN, Name, Salary, DNO)
- ☐ JobApplications(AppNo, ExpectedSalary, HireScore, Experience, SalaryOffered, Hired)
- ☐ Account(AccNo, Balance)
- ☐ Transaction(AcNo, TS, Description, CrDr, Amuont)

Concurrent execution of (SQL) statements

- Consider execution of following SQL statements from two concurrent clients (users).

C1: **UPDATE account SET balance = balance + 3000
WHERE accno = 1234;**

C2: **UPDATE account SET balance = balance - 5000
WHERE accno = 1234;**

- DBMS might execute them in interleaved fashion in following order –
 - C1 reads balance (10000)
 - C2 reads balance (10000)
 - C1 modifies it to (13000),
 - C2 modifies to 5000, C2 writes to database 5000,
 - C1 writes to database 13000

Issues with concurrent execution : Case-1

C1: **UPDATE account SET balance = balance + 3000
WHERE accno = 1234;**

C2: **UPDATE account SET balance = balance - 5000
WHERE accno = 1234;**

- ❑ DBMS executes them in following order –
 - C1 reads balance (10000)
 - C2 reads balance (10000)
 - C1 modifies it to (13000),
 - C2 modifies to 5000, C2 writes to database 5000,
 - C1 writes to database 13000
- ❑ Assuming that before both statements begin their execution, database has balance of 10000 for account no 1234.
- ❑ What will be and what should be final value of balance?
Do you see the problem?

Issues with concurrent execution : Case-2

C1: **UPDATE employee SET salary=60000
WHERE ssn = 1234;**

C2: **UPDATE employee SET dno=4 WHERE ssn = 1234;**

- ❑ DBMS may not be reading and writing only a single value from data on disk. Consider if smallest size DBMS reads and write is a tuple.
- ❑ Consider if above two statements are executed in following order –
 - C1 reads the tuple (ssn=1234)
 - C1 modifies the tuple, sets salary to 60000
 - C2 reads the tuple (ssn=1234)
 - C2 modifies the tuple, sets dno=4
 - C1 writes the tuple to the database
 - C2 writes the tuple to the database
- ❑ Note the problem?

Issues with concurrent execution : Case-3

C1: UPDATE JobApps SET OFFER = 40000
WHERE AppNo IN (SELECT AppNo FROM JobApps WHERE
score > 65);

C2: UPDATE JobApps SET score = 1.1*score
WHERE exp >= 5;

- C1: for each job_app x from JobApps
 read tuple x;
 If x.score > 65 then
 x.offer = 40000;
 write x;
- C2: for each job_app x from JobApps
 read tuple x;
 x.score += 1.1*x.score;
 write x;

Issues with concurrent execution : Case-3

C1: **UPDATE salary SET salary = 1.1*salary;**

C2:

SELECT avg(salary) FROM employee WHERE ;

SELECT avg(salary) FROM employee;

□ C1: **for each employee e from EMP**
 read tuple e;
 e.salary = 1.1 * e.salary;
 write e;

□ C2: **s=0;n=0;**
 for each employee e from EMP
 read tuple e;
 s += e.salary;
 write e;

Issues with concurrent execution : Case-4

C1: **INSERT INTO ex_emp** **SELECT .. FROM emp**
WHERE ex_flag;

DELETE FROM emp WHERE ex_flag;

C2: **SELECT count(*) FROM emp;**

SELECT count(*) FROM ex_emp;

☐ Note the problem?

Issue - System Failures

- Suppose, we need to transfer amount of 5000 from account number 1011 to 2312; and we have following database updates to log the transaction and update the balances accordingly.

```
INSERT INTO transaction VALUES (1011, now,  
    'Transferred to A/c 2312','Debit',5000);  
INSERT INTO transaction VALUES (2312, now,  
    'Transferred from A/c 1011','Credit', 5000);  
UPDATE account SET balance=balance-5000 WHERE  
    acno = 1011;  
UPDATE account SET balance=balance+5000 WHERE  
    acno = 2312;
```

- What if system crashes after executing 3rd statement ? We ought to execute either all statements or none ?

Conclusion drawn from issues

- ☐ Uncontrolled sequencing of operations for execution may bring database in inconsistent state!!
- ☐ Partial execution of a set of statements may bring database in inconsistent state.

Solution

- Control the concurrency
 - Order the Read/Write operations from multiple clients such that they appear to be running in isolation
- Guarantee executing sequence of related operations (Transactions) from a client in “all or nothing” manner, regardless of failures
- “Transaction” are to address both the above concerns

Notion of “Transaction”

- ❑ A transaction is a sequence of one or more SQL operations treated as a unit
- ❑ A sequence of SQL operations, performing a single application task; for example
 - Transfer amount from one account to another account; a sequence of related sql statements form a transaction
 - Save an Invoice; a sequences of statements that typically makes entry into multiple invoice tables; and update stock tables form a Transaction

SQL commands related to Transactions

- ❑ BEGIN TRANSACTION
- ❑ COMMIT
- ❑ ROLLBACK
- ❑ SET TRANSACTION
 - ISOLATION LEVEL
 - ACCESS MODE

ACID properties of Transaction

- ACID is acronym for
 - Atomicity
 - Consistency
 - Insolation
 - Durability

- These are desirable characteristics of Transaction Processing

ACID - Isolation

- ❑ If we execute transaction in isolation, that is no interleaving, DBMS let finish one transaction before taking up turn of another.
- ❑ That means no interleaving; that we can not afford – wait time, processor under-load, infinite loops in a transaction etc., etc.
- ❑ This conflicting situation brings in notion of **Serializability**.

ACID – Isolation (Serializability)

- ❑ Serializability: Operations may be interleaved, but execution must be equivalent to *some* **serial** (sequential) **order** of all transactions
- ❑ That is suppose there are two transactions T1 and T2 are executing concurrently; interleaving of operations from T1 and T2 should be such that the result is equivalent to either T1 followed by T2 executes in isolation, or T2 followed by T1.
- ❑ Notationally we say that “**T1;T2**” or “**T2;T1**”.
- ❑ So basically Isolation property requires transaction execution to ensure serializability.

ACID – Isolation (Serializability)

- ❑ So basically Isolation property requires transaction execution to ensure *Serializability*.
- ❑ With this promise, we can see that issues related to concurrency discussed earlier all gone.
- ❑ Case 1:
C1: **UPDATE account SET balance = balance + 3000
WHERE accno = 1234;**
C2: **UPDATE account SET balance = balance - 5000
WHERE accno = 1234;**
- ❑ No issue when executed as either C1;C2 or C2;C1

ACID – Isolation (Serializability)

- In case-2, and case-3 also, serializability ensures correct result.

Case-2

C1: **UPDATE employee SET salary=60000
WHERE ssn = 1234;**

C2: **UPDATE employee SET dno=4 WHERE ssn = 1234;**

ACID – Isolation (Serializability)

- ❑ However in case-3, different order produces different result, and might be desirable to other.
- ❑ Serializability does not guarantee the order of execution, it has to be controlled by the programmer

Case-3

C1: **UPDATE JobApps SET OFFER = 40000
WHERE AppNo IN (SELECT AppNo FROM JobApps WHERE
score > 65);**

C2: **UPDATE JobApps SET score = 1.1*score
WHERE exp >= 5;**

ACID – Isolation (Serializability)

- Case 4 also produces correct result; and this case also order will matter the result.

C1: **INSERT INTO ex_emps SELECT .. FROM emp WHERE ex_flag;**

DELETE FROM emp WHERE ex_flag;

C2: **SELECT count(*) FROM emps;**

SELECT count(*) FROM ex_emps;

ACID - Durability

- ❑ This characteristic requires that if system crashes after transaction commits; all effects of transaction reflect in database.
- ❑ In real systems, DBMS might acknowledge the client for a commit request, and a system failure occurs before effects of committed transaction are updated on disk.
- ❑ A transaction processing system should prevent this and ensure --> Durability.

ACID - Atomicity

- ❑ A transaction is executed in “all or nothing” manner.
- ❑ A transaction is never half executed.
- ❑ DBMS ensures this by maintaining appropriate logs, and by recovery process.

ACID - Consistency

- ❑ Each transaction from each client can assume that all constraints hold when transaction begins.
- ❑ Transaction execution should ensure that all constraints hold true after transaction commits.
- ❑ Serializability and atomicity takes care of this property.

Repeat

- Study of Transactions addresses following two requirements –
 - Enable concurrent access of databases - multiple users need to concurrently access and update databases.
 - Resilience to system failures - While databases are updated system might crash in between.

□ SQL Transaction Isolation Levels

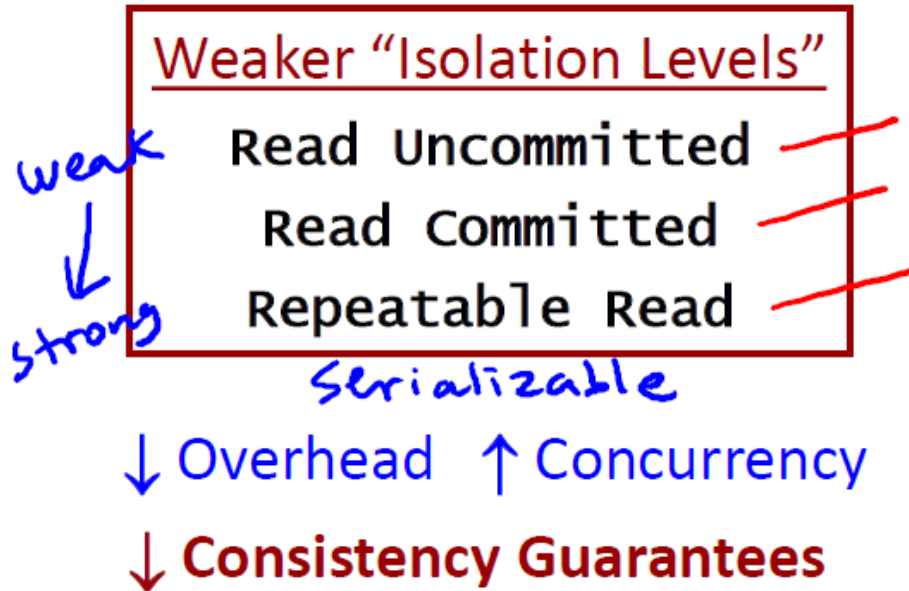
SQL Isolation Levels

- ❑ While serializability is desirable; it has following side effects -
 - Brings in additional overhead (processes that ensure serializable execution schedules)
 - Provides reduced concurrency; Serialization algorithms may ask some transaction to wait and aborts
- ❑ However Serializability may not required all the time, particularly for reads; therefore SQL allows specifying a transactions to execute “weaker isolations”.

SQL Transaction Isolation Levels

- SQL standard specifies four Isolation levels in which a transactions can be executed

- Serializable
- Repeatable Read
- Read Committed
- Read UnCommitted



Source: Lecture slides from Jennifer Widom

SQL Transaction Isolation Levels

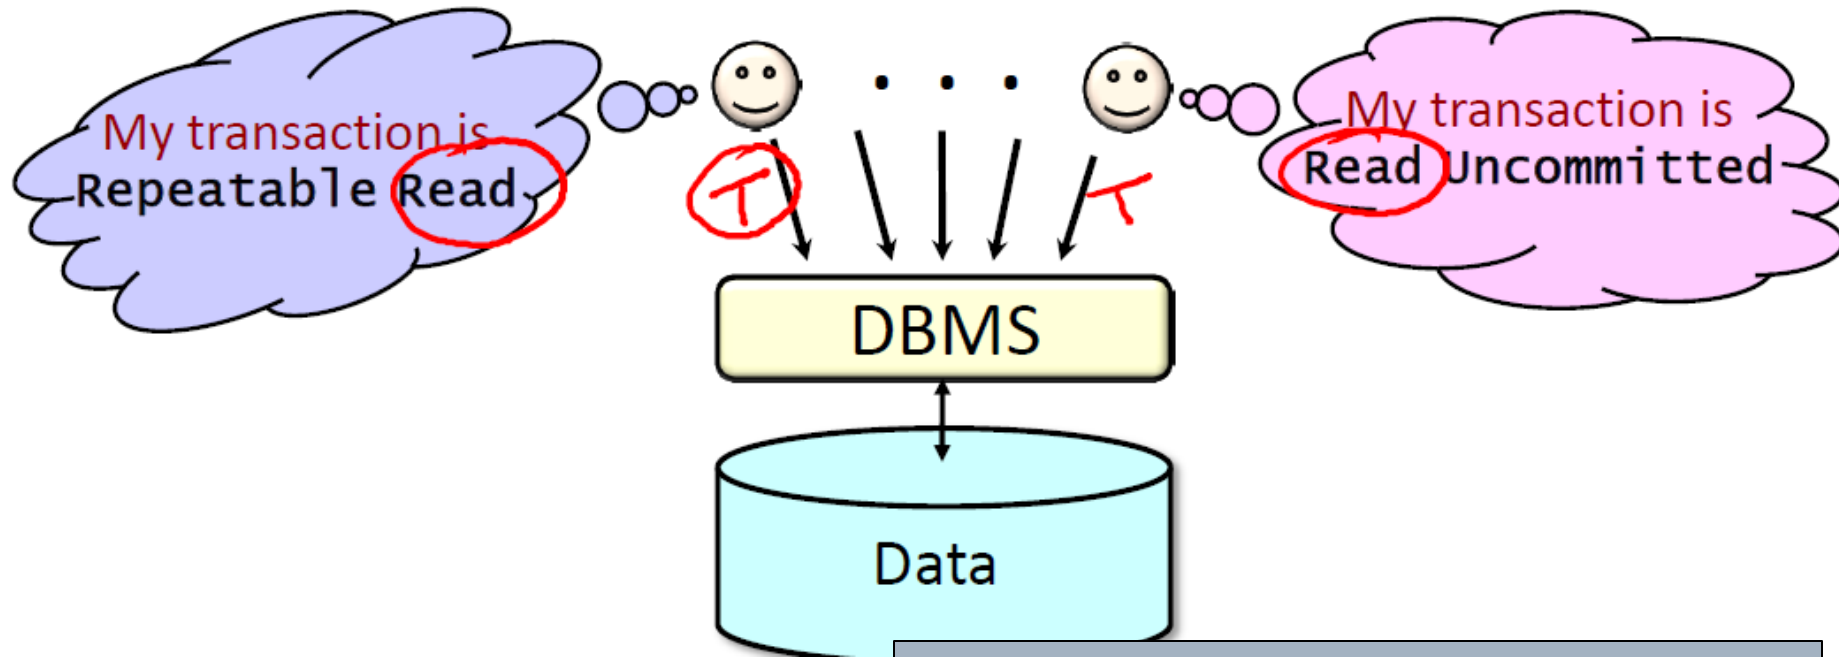
- ❑ Specified for the transaction
- ❑ Different transaction may execute at different isolation levels
- ❑ SQL has SET Transaction command to specify isolation level, and done as following -

`Set Transaction Isolation Level <isolation level>;`

SQL Transaction Isolation Levels

Isolation Levels

- Per transaction
- “In the eye of the beholder”



Source: Lecture slides from Jennifer Widom

Transaction isolation levels in SQL-92

☐ **Serializable**

☐ **Repeatable read** —

- only committed records to be read,
- repeated reads of same record must return same value; however, a transaction may not be serializable- it may find some records inserted by other concurrent transactions.

☐ **Read committed** —

- only committed records can be read,
- but successive reads of same record may return different (but committed only) values.

Transaction isolation levels in SQL-92

- ❑ **Read uncommitted** –
 - Even uncommitted records may be read.
 - Dirty read may occur
- ❑ As per SQL standard, **Serializable** is the default isolation level.
- ❑ However a RDBMS may have different level as default. For example PostgreSQL has **read committed** as default isolation level.

☐ Problems with lower isolation levels

Isolation Level – Read Uncommitted

T1:

```
UPDATE emp SET salary = 1.1*salary;  
ROLLBACK; //COMMIT
```

T2: **SELECT average(salary) FROM emp;**

- ❑ Suppose T2 runs at “Read UnCommitted” level, then it allows reading data from T1 even before it commits. If T1 decides to abort then report of T2 is incorrect.
- ❑ Even when T1 commits, T2 might read some old data before T1 updates and some updated that also undesirable.
- ❑ This phenomena of reading uncommitted data is called as “Dirty Read”.

Isolation Level – Read Uncommitted

- A transaction running at **Read Uncommitted** isolation level may perform dirty reads
- Note that in previous example if T2 performs dirty then it is not ensuring serializability; that is

Neither “T1;T2” nor “T2;T1”

Isolation Level – Read Committed

- ❑ A transaction running at **Read Committed** level does not read updates from uncommitted transaction.
- ❑ Consider concurrent execution of following transaction
T1: **UPDATE emp SET salary = 1.1*salary;**
COMMIT;
T2: **SELECT average(salary) FROM emp;**
SELECT max(salary) FROM emp;
- ❑ If first statement of T2 executes (in Read Committed level) while T1 executes its UPDATE; it reads from old data.
- ❑ And by the time T2 starts executing second statement, T1 commits, so second statements reads data updated by T1

Isolation Level – Read Committed

T1: **UPDATE emp SET salary = 1.1*salary;
COMMIT;**

T2: **SELECT average(salary) FROM emp;
SELECT max(salary) FROM emp;**

- ☐ The said execution is neither T2;T1 nor T1;T2.
- ☐ Execution of transaction at this level does not perform dirty read, still does not guarantee serializability.
- ☐ It is due to insufficient isolation level

Isolation Level – Repeatable Read

- A transaction reads same value if a data item is read multiple times; that is repeated read of a item gets same value in a transaction.
- In previous example, problem was due to non repeatable reads. Now if T2 in following concurrent execution runs at “Repeatable Read; read of salary in both statements of T2 will be same ==> results serializable execution

T1: **UPDATE emp SET salary = 1.1*salary;**
COMMIT;

T2: **SELECT average(salary) FROM emp;**
SELECT max(salary) FROM emp;

Isolation Level – Repeatable Read

- Consider concurrent execution of following transactions-

T1: **INSERT INTO emp .. ;** //say inserts 10 tuples
COMMIT;

T2: **SELECT average(salary) FROM emp;**
SELECT average(salary) FROM emp;

- If first statement of T2 executes (in Repeatable Read) while T1 executes its UPDATE; it reads from old data.
- And by the time T2 starts executing second statement, T1 commits, so second statements reads data updated by T1
- This execution again will not achieve serializability?
- Repeatable reads sees rows inserted by other concurrent transactions, called “phantom rows”

Isolation Level – Repeatable Read

□ Consider concurrent execution of following transactions-

T1: **UPDATE emp SET salary = 1.1*salary;**

UPDATE emp SET dno=4 WHERE ssn=1234

COMMIT;

T2: **SELECT average(salary) FROM emp;**

SELECT count(*) FROM emp WHERE dno=4;

- If first statement of T2 executes (in Repeatable Read) while T1 executes its UPDATE; it reads from old data.
- And by the time T2 starts executing second statement, T1 commits, so second statements reads data updated by T1
- This execution may not achieve serializability? If concurrent transaction

Isolation Level - Summary

wrk

| | <i>←</i> dirty reads | <i>←</i> nonrepeatable reads | <i>←</i> phantoms |
|------------------|-------------------------|------------------------------------|----------------------|
| Read Uncommitted | Y | Y | Y |
| Read Committed | N | Y | Y |
| Repeatable Read | N | N | Y |
| Serializable | N | N | N |

Strong

Source: Lecture slides from Jennifer Widom

Isolation Level - Summary

- ❑ Standard default: Serializable
- ❑ Weaker isolation levels
 - Increased concurrency + decreased overhead = increased performance
 - Weaker consistency guarantees
 - Some systems have default Repeatable Read
- ❑ Isolation level per transaction and “eye of the beholder”
 - Each transaction’s reads must conform to its isolation level

Source: Lecture slides from Jennifer Widom

□ Source: Lecture slides from Jennifer Widom